

# Fitting a Forth in 512 bytes

June 10, 2021 · 39 minute read

This article is part of the <u>Bootstrapping</u> series, in which I start from a 512-byte seed and try to bootstrap a practical system.

#### Next post:

No branches? No problem — a Forth assembler

Software is full of circular dependencies if you look deep enough. Compilers written in the language they compile are the most obvious example, but not the only one. To compile a kernel, you need a running kernel. Linkers, build systems, shells. Even text editors, if you want to write the code instead of just downloading it. How do you break this cycle? Since the bootstrapping problem has first come to my attention, I've been drawn to this unique area of software engineering. Not out of fear that someone would try to implement a trusting trust attack, but simply as an interesting challenge.

11 years ago, <u>vanjos72</u> described on Reddit what he calls a thought experiment: what if you were locked in a room with an IBM PC, with no operating system on it? What would be the minimum amount of software you'd need to start out with to bootstrap back into comfort?

As it happens, I've recently found myself with an abundance of free time on my hands, so I've decided to make this more than a thought experiment. Alas, my computer didn't come equipped with front panel switches, so some software needs to be present on the computer already...

The absolutely minimal option would be a simple program that accepts input from the keyboard, and then jumps to it. Since the keyboard input routines in the BIOS implement alt+numpad escape codes, you don't even need to write any base conversion code. Moreover, the loop doesn't even need an end condition — just write to the buffer backwards until you run into the existing code and overwrite the jump target. This approach takes a mere 14 bytes:

https://niedzejkob.p4.team/bootstrap/miniforth/

```
aa stosb
ebf9 jmp short input_loop
buffer:
```

However, I do not find the prospect of entering code this way anywhere near appealing. I've decided that, since the BIOS loads an entire sector anyway, any bootstrap seed that fits into the bootsector is fair game.<sup>3</sup> Obviously, one would want to maximize the utility of the chosen program. What is the most powerful thing we can fit in 510 bytes?

Many interesting sector-sized programs have been written by Oscar Toledo. This includes many games, such as a <u>DooM-like raycasting game</u> or a <u>chess AI</u>, as well as a basic <u>BASIC interpreter</u>, but the most perhaps the most relevant one for our usecase is <u>bootOS</u>:

bootos is a monolithic operating system that fits in one boot sector. It's able to load, execute, and save programs. Also keeps a filesystem.

It exposes its filesystem routines with an interrupt interface, and includes a builtin command that allows creating a file by typing in its hexdump. Very neat, but clearly mostly intended as a multiplexer between other sector-sized programs.

What I would seek is a solution that minimizes typing in hand-assembled machine code. Ideally, it would be a programming language, but one that, unlike BASIC, can be extended at runtime. If you've read the title of this post, you already know what I settled on — as it turns out, it's possible to fit a barebones Forth in a bootsector. You can see the code in the Miniforth repository on GitHub, but I will include most of it here.

The entire Forth takes, at this moment, 504 bytes. As you might expect, the development process involved being on a perpetual lookout for byte-saving opportunities. However, when I published what I thought was quite tightly optimized code, <u>Ilya Kurdyukov</u> came along and managed to find 24 bytes to be saved! I promptly reinvested this saved space in <u>new features</u>.

# A primer on Forth

If you've ever written anything in Forth, you can safely skip this section.

Forth is a stack-based language. For example, a number will push its value onto the stack, while the + word will pop two numbers and push their sum. A common debugging utility, but one not included in Miniforth, is the .s word, which prints the contents of the stack.

```
1 2 3 + .s <2> 1 5 ok
```

The user can define their own words with : and ; . For example:

```
: double dup + ; ok
3 double . 6 ok
```

This defines the word double, which does the same thing as dup + . dup, by the way, is one of Forth's stack manipulation words. It duplicates the top element on the stack:

```
42 dup .s <2> 42 42 ok
```

This is basically the entire language. There are some standard facilities for conditionals and loops, but we don't need to concern ourselves with those for now, as they can be built on top of Miniforth later on.

To talk about the effect a word has on the state of the stack, we use a notation like this:

```
dup ( a -- a a )
swap ( a b -- b a )
```

The list before the -- are the inputs, with the top of stack listed last. After the -- , we list the outputs, which start at the same stack depth. This lets us succintly describe the common aspects of a word.

#### Threaded code

While some Forth systems do include full-blown, optimizing compilers similar to those one'd see in a typical programming language, there is a much simpler strategy. After all, everything a Forth word can do is execute other words, so a sequence of call instructions gets us very close:

```
DOUBLE:

call DUP

call PLUS

ret
```

However, this ties up the hardware x86 stack for the return stack, making us handroll a separate stack for the actual user-level stack (known as the *parameter stack*). As accessing the parameter stack is much more common, we'd like to use the push and pop instructions for that, and instead handroll a mechanism similar to call. Firstly, let's simply store a list of pointers to words:

```
DOUBLE:

dw DUP

dw PLUS
```

The way this comes to life is that each primitive word fetches the address of the next word from memory, and jumps to it. A pointer to this sequence of pointers is kept in <code>si</code>, so that the <code>lodsw</code> instruction allows for easy processing of this list:

```
pop ax
push ax
push ax
lodsw
jmp ax

PLUS:
pop ax
pop bx
add ax, bx
push ax

lodsw
jmp ax
```

This common code can be abstracted away into a macro, which is traditionally called NEXT:

```
%macro NEXT 0
    lodsw
    jmp ax
%endmacro
```

This mechanism, by the way, is known as threaded code. No relation to the concurrency primitive.

What happens if one compiled word calls another one, though? This is where the return stack comes in. It might feel natural to use the BP register for this stack pointer. However, in 16-bit x86, there isn't actually a [bp] addressing mode. The closest you can get is [bp+imm8], which means that accessing the memory at bp wastes a byte to specify that you do not want an offset. This is why I use the di register for the return stack instead. Overall, this choice saves 4 bytes.

Anyway, here is how the return stack is used to handle compiled words calling each other. Pushing onto the return stack is particularly nice, since it's just the stosw instruction.

```
DOUBLE:

call DOCOL

dw DUP

dw PLUS
```

```
DOCOL: ; short for "do colon word"

xchg ax, si; used here as `mov ax, si`, but swaps with

; ax are only one byte, while `mov`s are two bytes

stosw

pop si; grab the pointer pushed by `call`

NEXT

EXIT:

dec di
dec di
mov si, [di]
NEXT
```

This is pretty much the execution strategy used by Miniforth, with one simple, but significant improvement — the value on top of the stack is stored in the BX register. This allows skipping a push and pop in many primitives:

```
PLUS:

pop ax

add bx, ax

NEXT

DROP:

pop bx

NEXT

DUP:

push bx

NEXT
```

One case is still unresolved, though. What happens if a word contains a number, such as : DOUBLE 2 \* ; ? This is handled by LIT, which will fetch the literal that follows out of the pointer stream:

```
DOUBLE:

call DOCOL

dw LIT, 2

dw MULT

dw EXIT

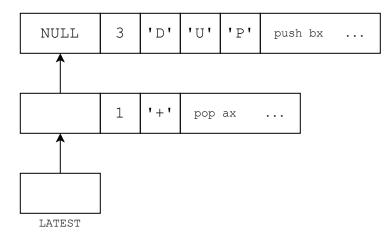
LIT:

push bx
```

```
lodsw xchg bx, ax
```

# The dictionary

We need a way to locate the implementation of the words the user types in. This is the role of the *dictionary*. I use a structure similar to many other small-scale Forths — a singly linked list of word headers, directly prepended before the code of each word. Out of tradition, the head of the list is kept in a variable called LATEST.



The most significant bits of the name length field also store some flags:

```
F_IMMEDIATE equ 0x80
F_HIDDEN equ 0x40
F_LENMASK equ 0x1f
```

If a word is marked as <code>IMMEDIATE</code>, it will be executed immediately, even if we're currently compiling a definition. For example, this is used to implement <code>;</code>. If a word is marked as <code>HIDDEN</code>, it is ignored when searching through the dictionary. Apart from being used as a rudimentary encapsulation mechanism, this can be used to implement the traditional Forth semantics where a word definition can refer to the previous word with the same name (and <code>RECURSE</code> is used when you want the definition currently being compiled). However, towards the end of development, I have removed the code that actually does this from the default implementation of <code>:</code> and <code>;</code>.

## Compression

It is usually not worth it to use compression when both the decompressor and its payload have to fit in merely 512 bytes. However, in a Forth implementation, one thing that's repeated very often is the implementation of  $_{
m NEXT}$ .

```
ad lodsw ffe0 jmp ax
```

We could try to save some bytes by replacing these with jumps to a shared copy. However, a short jump still takes two bytes — not a significant saving. As it turns out, a special compression scheme that can only handle this one repeating pattern is worth it, as long as you combine it with the following observation: NEXT is almost always followed by the dictionary entry of the next primitive, of which the link field is predictable.

I chose to implement a compression scheme where every  $0 \times ff$  byte is replaced with NEXT, followed by a link field, which is computed based on the previous occurence of an  $0 \times ff$  byte. This strategy saved 19 bytes when I introduced it.<sup>4</sup>

At first, I used a 0x90 byte for this — after all, it's the opcode of nop, which I'm definitely not going to be using. However, the byte can still occur in the immediate bytes of an instruction. It wasn't a problem at first, but when the code was shifting around in memory, various addresses and offsets became 0x90 often enough to be a nuisance. 0xff doesn't seem to have this problem.

To create a link, we copy the value of LATEST to the decompressor output, and update LATEST to point to the word we've just written. This can be done in a very compact sequence of instructions, but it still takes enough bytes that it is worthy it to factor it out as a subroutine — it is also used by the implementation of : , which creates dictionary entries at runtime.

The decompressor used to make use of an interesting trick, where instead of a short forward jump, an opcode is placed so the immediate argument it requires eats the instructions we want to jump over. That is, instead of

```
jmp short .after
.write:
    stosb
.after:
```

you write

```
3c db 0 \times 3c; skip the stosb below by comparing its opcode with AL .write: aa stosb
```

Thus, if some other code jumps to .write , the stosb executes, but this codepath just does cmp al, 0xaa . At first, I didn't think of the cmp al instruction, and a mov into a throwaway register instead. This backfired spectacularily because of my inability to actually pick a register that can be safely overwritten.

Ilya Kurdyukov then demonstrated that the same bytecount can be achieved without this kind of "magic". An analogous modification allowed me to remove the other occurence of this trick too. The essence is that instead of trying to skip over the stosb, we execute it unconditionally before the codepaths branch, and then essentially undo it with dec di if necessary:

```
SPECIAL BYTE equ 0xff
   mov si, CompressedData
   mov di, CompressedBegin
   mov cx, COMPRESSED SIZE
.decompress:
   lodsb
   stosb
   cmp al, SPECIAL BYTE
   jnz short .not special
   dec di
   mov ax, 0xffad ; lodsw / jmp ax
   stosw
   mov al, 0xe0
   stosb
   call MakeLink
.not special:
   loop .decompress
```

Actually generating the compressed stream is more involved. Because I want jumps between the compressed and uncompressed portions to work, the assembler needs to believe it is writing the code at the location it will actually run. I first attempted to do this by adjusting the <code>org</code> after each <code>SPECIAL\_BYTE</code>, but unfortunately, yasm didn't like that.

```
boot.s:137: error: program origin redefined
```

Clearly, a separate post-processing step is necessary. I wrote a macro to shim the bytes the decompressor will insert:

```
%macro compression_sentinel 0
    db SPECIAL_BYTE
    dd 0xdeadbeef
%endmacro
```

This has the added benefit of allowing a simple automated way to verify that no SPECIAL\_BYTE s slipped in by accident.

I still had to allocate the space for the compressed data. I choose the following layout:

- 1. Uncompressed code starts at 7000 initialization, decompression, and the outer interpreter.
- 2. Compressed data immediately follows, filling up the bootsector up to a moment before 7E00.
- 3. The decompression buffer is allocated immediately after that, which is where yasm outputs the target contents.

To achieve this, I needed to know exactly how much space needs to be allocated for the compressed data. First, I calculate the exact number of bytes saved by incrementing a counter in the compression\_sentinel macro:

```
%assign savings 0

%macro compression_sentinel 0
%assign savings savings+4
    db SPECIAL_BYTE
    dd 0xdeadbeef
%endmacro
```

Then, I simply subtract this from the size of the uncompressed segment:

```
CompressedData:
    times COMPRESSED_SIZE db 0xcc

CompressedBegin:
; ...
CompressedEnd:

COMPRESSED_SIZE equ CompressedEnd - CompressedBegin - savings
```

The post-processing is done by a simple Python script:

```
SPECIAL_BYTE = b'\xff'
SENTINEL = SPECIAL_BYTE + b'\xef\xbe\xad\xde'
```

```
with open('raw.bin', 'rb') as f:
    data = f.read()
# Recognize the reserved space by an arbitrary, but relatively large threshold
# of 20 repeated \xcc bytes.
output offset = data.index(b'\xcc' * 20)
chunks = data[output offset:].lstrip(b'\xcc').split(SENTINEL)
assert SPECIAL_BYTE not in chunks[0]
compressed = bytearray(chunks[0])
for chunk in chunks[1:]:
    assert SPECIAL BYTE not in chunk
    compressed.extend(SPECIAL BYTE)
    compressed.extend(chunk)
# Make sure that exactly the right amount of space is allocated
# for the compressed data.
assert b'\xcc' * len(compressed) in data
assert b'\xcc' * (len(compressed) + 1) not in data
output = data[:output offset] + compressed
print(len(output), 'bytes used')
output += b' \times 00' * (510 - len(output))
output += b' \times 55 \times aa'
with open('boot.bin', 'wb') as f:
    f.write(output)
```

The same script also generates an extended disk image, which contains some smoke-testing code in block 1:

```
output += b'\x00' * 512
output += open('test.fth', 'rb').read().replace(b'\n', b' ')
output += b' ' * (2048 - len(output))
with open('test.img', 'wb') as f:
    f.write(output)
```

compression\_sentinel is most often used by the defcode macro, which creates the dictionary entry for a primitive word. It takes a label (which can then be used to jump to the implementation of some word), the name of the word as a string, and optionally, some flags to be ORed into the length

field:

```
; defcode PLUS, "+"
; defcode SEMI, ";", F_IMMEDIATE
%macro defcode 2-3 0
    compression_sentinel
%strlen namelength %2
    db %3 | namelength, %2
%1:
%endmacro
```

This is then used to define the primitives. The code essentially falls-through into a defcode:

```
defcode PLUS, "+"
   pop ax
   add bx, ax

defcode MINUS, "-"
   pop ax
   sub ax, bx
   xchg bx, ax

defcode PEEK, "@"
   ; ...
```

However, DOCOL, EXIT and LIT also use the compression mechanism for their NEXT s. Since the link field is still written out, this essentially creates bogus dictionary entries. Fortunately, the first opcode of EXIT and LIT has the F\_HIDDEN bit set, so this is not a problem:

```
DOCOL:
    xchg ax, si
    stosw
    pop si; grab the pointer pushed by `call`
    compression_sentinel

LIT:
    push bx
    lodsw
    xchg bx, ax
    compression_sentinel

EXIT:
```

```
dec di
  dec di
  mov si, [di]

defcode PLUS, "+"
; ...
```

#### Variables?

Immediate load instructions tend to have shorter encodings than loads from memory:

```
be3412 mov si, 0x1234
8b363412 mov si, [0x1234]
```

This is why Miniforth stores most of its variables in the immediate fields of instructions (a practice known as self-modifying code). Of course, this means that the address of these variables will change on every edit of the code, which is problematic, since we will be wanting to access these variables in Forth code. The typical way of exposing a variable is to create a word that pushes its address. However, that's way too expensive with our constraints. What I settled on is pushing the addresses onto the stack at startup. This can be done with only 2 bytes for each address, by simply defining the initial contents of the stack as data:

```
jmp 0:start
stack:
   dw HERE
   dw BASE
   dw STATE
   dw LATEST
start:
   ; ...
   mov sp, stack
```

Even when a variable's address needs to be pushed onto the stack, this self-modifying code strategy saves bytes if a variable needs to be initialized — the best way to initialize a variable is to simply allocate it within the bootsector and dw the initial value there, which exactly evens out the stack data, and keeps the advantage of the shorter instruction encoding.

#### Initialization code

The first thing done after booting is setting up the <u>segment registers</u> and stack. The direction flag is

also cleared, so that the string instructions work in the right direction.

```
jmp 0:start
; ...
start:
   push cs
   push cs
   push cs
   pop ds
   pop es
   pop es
   pop ss
   mov sp, stack
cld
```

There are a two notable things about this code. Firstly, segment registers are set through the stack. This is a byte-saving trick I've picked up from <code>bootbasic</code> — it allows having to initialize a general-purpose register to zero:

```
31c0
      xor ax, ax ; through AX - 8 bytes
8ed8 mov ds, ax
8ec0
     mov es, ax
8ed0
       mov ss, ax
0e
                ; through the stack - 6 bytes
       push cs
0e
       push cs
0e
       push cs
1f
       pop ds
07
       pop es
17
       pop ss
```

Secondly, one would think that, while the stack is being repointed, a small race condition window occurs — if an interrupt happened between <code>pop ss</code> and <code>mov sp</code>, chaos could ensue if the previous value of SP was in an unlucky place in memory. Of course, I could just cross my fingers and hope this doesn't happen if the 2 bytes required to wrap this in an <code>cli/sti</code> pair were too much. However, it turns out that this trade-off is not necessary due to an obscure corner of the x86 architecture. To quote the Volume 2B of the x86 Software Developer's Manual:

Loading the SS register with a POP instruction<sup>5</sup> suppresses or inhibits some debug exceptions and inhibits interrupts on the following instruction boundary. (The inhibition ends after delivery of an exception or the execution of the next instruction.) This behavior allows a stack pointer to be loaded into the ESP register with the next instruction (POP ESP),<sup>6</sup> before an event can be delivered.

After the segments, stack and direction flag are set up, the decompressor is ran. Crucially, it does not use the DL register, which contains the BIOS disk number from which we were booted. It is then poked into the implementation of <code>load</code> (which is in the compressed segment), and pushed onto the stack for later use by user code:

```
mov [DRIVE_NUMBER], d1
push dx; for Forth code
```

### The outer interpreter

At this point, we reach the *outer interpreter* - the part of a Forth system that processes user input. The name "*outer* interpreter" distinguishes it from the *inner* interpreter, which is the component that coordinates the execution within a defined word, and consists of NEXT, DOCOL, EXIT, and LIT.

Normally, a Forth would expose the building blocks of its outer interpreter as words in the dictionary, such as

- REFILL (read a line of input from the currently executing source),
- WORD (parse a word from the input stream),
- FIND (look up a word in the dictionary),
- >NUMBER (convert a string to number).

In Miniforth, no attention is paid to this practice at all. Dictionary headers cost bytes, and so does communicating only through the stack. In fact, word and >NUMBER are melded together into one routine that does the job of both — that way, the loop can be shared, which saves bytes.

This monolithic architecture also lets us decide that BX and DI are not reserved for the top of stack and the return stack pointer, respectively, while the outer interpreter is executing. This significantly helps with register starvation within these comparatively complex parts of the system. These registers are set up just before jumping to a word, and saved after it returns.

### Keyboard input

After initialization is completed, the code falls through to ReadLine, the routine for reading in an input line from the keyboard. We will also jump back here later, when the current line of input is exhausted. The input buffer is at 0x500, directly after the BDA. While the idiomatic string format for Forth uses a separate length field, this buffer is NULL-terminated, as that is easier to handle when parsing. The pointer to the unparsed fragment of the input is stored in InputPtr, which is the only variable which does *not* use the self-modification technique, as it does not need to be explicitly initialized — it naturally gets written to before it is read.

```
InputBuf equ 0x500
InputPtr equ 0xa02 ; dw
   mov di, InputBuf
   mov [InputPtr], di
.loop:
   mov ah, 0
   int 0x16
   cmp al, 0x0d
   je short .enter
   stosb
   cmp al, 0x08
   jne short .write
   dec di
   cmp di, InputBuf; underflow check
   je short .loop
   dec di
.write:
   call PutChar
   jmp short .loop
.enter:
   call PutChar
   mov al, 0x0a
   int 0x10
   xchg ax, bx; write the null terminator by using the BX = 0 from PutChar
   stosb
InterpreterLoop:
   call ParseWord; returns length in CX. Zero implies no more input.
   jcxz short ReadLine
```

The BIOS interrupt for getting a character from the keyboard does not print the key — we have to do that ourselves. This is done with the "TELETYPE OUTPUT" function, which already handles special characters like backspace or newline.

```
PutChar:

xor bx, bx

mov ah, 0x0e

int 0x10

ret
```

This function has its deficiencies. For example, the icky CRLF line endings are needed (CR to move the cursor to the beginning of the line, and LF to move it to the next line). Also, the backspace character only moves the cursor back a character, and does not erase it. To get the behavior we've

come to expect, it would be necessary to print \b \b (to be fair, this is also the case on modern terminals). I chose to skip that.

Finally, <u>Ralf Brown's Interrupt List</u> mentions that some BIOSes clobber BP when the printed character causes the screen to scroll. This does not concern us, as we do not use this register at all.

# Parsing

After we read in a line, we need to parse it into words. This is done on demand — each word is executed (or compiled, depending on the state), as soon as it is parsed. Apart from the interpreter loop, Parseword is also called by the implementation of : (to get the name of the word being defined).

As mentioned before, this routine also computes the numeric value of the word, with the assumption that it's valid. There is no error checking in this regard — if a word is not found in the dictionary, its numeric value is pushed, which is probably nonsense if this wasn't intended.

We start out by skipping any whitespace in the input buffer:

Note the way the pointer to the beginning of the string is saved. The loop will go one byte past the whitespace, so storing it after the loop would require a separate decrement. Instead, we update the register with each iteration of the loop, but before the pointer is incremented by lodsb.

At this point, the AL register is loaded with the first character of the word. Thus our next loop will need to do lodsb at its *end*.

```
xor cx, cx
```

```
xor bx, bx
.takeloop:
   and al, ~0x20
   jz short Return ; jump to a borrowed `ret` from some other routine
```

This and instruction is interesting, as it does three things at once. It detects both spaces and null bytes in one fell swoop, but also also turns off the bit that differs between uppercase and lowercase letters, which allows handling both cases of hexadecimal numbers at no extra cost.

If we haven't detected the end of the word, we increment the length counter and convert the digit to its numeric value:

```
inc cx
sub al, "0" &~0x20
cmp al, 9
jbe .digit_ok
sub al, "A" - ("0" &~0x20) - 10
.digit_ok
cbw
```

cbw is a little-known instruction that converts a signed number from b yte to w ord, but for us it's just a shorter mov ah, 0. In a perhaps similar vein, we use the signed multiply imul, because it has more options for how it uses the registers than the unsigned mul. The particular form used here allows multiplying by an immediate and doesn't overwrite DX with the upper half of the product.  $^{7}$ 

This particular instruction needs to be encoded manually to force the literal width to be 2 bytes wide.<sup>8</sup>

```
; imul bx, bx, <BASE> but yasm insists on encoding the immediate in just one byte. db 0x69, 0xdb BASE equ $ dw 16 add bx, ax; add the new digit
```

Finally, we set up for the next iteration of the loop. We use a similar trick as before, where a pointer result is updated in each iteration to avoid a separate decrement at the end — we need to make sure that the input pointer doesn't point after the terminator.

```
mov [InputPtr], si
lodsb
jmp short .takeloop
```

# **Dictionary Lookup**

After a word is parsed, we try to look it up in the dictionary. For each entry, we need to compare the length of the name, and if it matches, the name itself. By including <code>F\_HIDDEN</code> in the mask, we automatically handle hidden entries, too. The way we're comparing the length might look a bit weird. The goal is to keep the <code>F\_IMMEDIATE</code> bit in AL, so that we don't have to keep around the pointer to the header of this word. This is one of Ilya Kurdyukov's clever optimizations.

```
InterpreterLoop:
   call ParseWord
    jcxz short ReadLine
; Try to find the word in the dictionary.
; SI = dictionary pointer
; DX = string pointer
; CX = string length
; Take care to preserve BX, which holds the numeric value.
LATEST equ $+1
    mov si, 0
.find:
    lodsw
    push ax ; save pointer to next entry
    xor al, cl; if the length matches, then AL contains only the flags
    test al, F HIDDEN | F LENMASK
    jnz short .next
    mov di, dx
    push cx
    repe cmpsb
    pop cx
    je short .found
.next:
    pop si
    or si, si
    jnz short .find
    ; If we reach this point, it's a number.
    ; ...
```

This part shows another advantage of not splitting the interpreter into reusable chunks — we can easily exit into two different codepaths, based on the result of the lookup.

### Should we execute it?

The system can be in two possible states:

- interpreting all words should be executed
- compiling immediate words should be executed

In other words, a word should be executed if it is immediate, or we're interpreting. We store this flag in the immediate field of an or instruction — it will be set to 0 when compiling:

```
; When we get here, SI points to the code of the word, and AL contains
; the F_IMMEDIATE flag

STATE equ $+1
  or al, 1
  xchg ax, si; both codepaths need the pointer to be in AX
  jz short .compile

; Execute the word
; ...
```

The most important words that need to change the state are : and ; , but they just jump to [ and ] — words that allow to temporarily drop back to interpreted mode while compiling a word. The typical usecase is to eagerly calculate the value of a constant expression:

```
: third-foo [ foos 3 cells + ] literal @ ;
```

Since the two values of STATE differ only by 1, we can switch between them with inc and dec. This has the disadvantage that they are no longer idempotent, but this shouldn't matter to well-written code:

```
defcode LBRACK, "[", F_IMMEDIATE
    inc byte[STATE]

defcode RBRACK, "]"
```

```
dec byte[STATE]
```

# Executing the word

If we decided to execute the word, we retrieve BX and DI, and set up SI so that NEXT will jump back to .executed:

```
; Execute the word
RetSP equ $+1
  mov di, RS0
  pop bx
  mov si, .return
  jmp ax
.return:
  dw .executed
.executed:
  mov [RetSP], di
  push bx
  jmp short InterpreterLoop
```

# Handling numbers

There is no F\_IMMEDIATE flag for numbers, so we just need to check the state to decide. It's a simple comparison, but if we're clever enough, you can save a byte here. Let's look again at the code that searches the dictionary. What value will AH have when we reach the number case?

```
.find:
   lodsw
   push ax ; save pointer to next entry
   lodsb
   xor al, cl ; if the length matches, then AL contains only the flags
   test al, F_HIDDEN | F_LENMASK
   jnz short .next

mov di, dx
   push cx
   repe cmpsb
   pop cx
   je short .found
.next:
   pop si
   or si, si
   jnz short .find
```

```
; AH = ?
```

Do you see it? At this point, AH is zero, since it contains the higher half of the pointer to the next word, which we know is NULL, as we just got to the end of the list. This allows us to check the value of STATE without loading it into a register or any immediate bytes:

```
; It's a number. Push its value - we'll pop it later if it turns out we need to cc
; it instead.
push bx
cmp byte[STATE], ah
jnz short InterpreterLoop
; Otherwise, compile the literal.
; ...
```

# Compiling things

The output pointer for the compilation process is called  $\mbox{HERE}$ . It starts out just after the decompressed data. The function that writes out a word into this area is called  $\mbox{comma}$ , since the Forth word that does this is  $\mbox{,}$ .

```
COMMA:

HERE equ $+1

mov [CompressedEnd], ax

add word[HERE], 2

ret
```

It is used in a straight-forward way to compile both numbers and words. We can share the tail, though — compiling a word will jump into the middle of compiling a number:

```
; Otherwise, compile the literal.
mov ax, LIT
call COMMA
pop ax
.compile:
    call COMMA
    jmp short InterpreterLoop
```

The last piece of the puzzle are : and ; . Let's look at : first. Since Parseword makes use of BX and SI, we need to save these registers. Moreover, since we're writing the many parts of a dictionary header, we'll load HERE to DI to streamline things. This is a lot of registers that we need

to push. However, we don't actually need to modify any register, so we can just save *all* the registers with <code>pusha</code>.

```
defcode COLON, ":"
   pusha
   mov di, [HERE]
   call MakeLink ; link field
   call ParseWord
   mov ax, cx
   stosb ; length field
   mov si, dx
   rep movsb ; name field
   mov al, 0xe8 ; call
   ; The offset is defined as (call target) - (ip after the call instruction)
   ; That works out to DOCOL - (di + 2) = DOCOL - 2 - di
   mov ax, DOCOL - 2
   sub ax, di
   stosw
   mov [HERE], di
   jmp short RBRACK ; enter compilation mode
```

; is much shorter. We merely need to compile EXIT and go back to interpretation mode:

```
defcode SEMI, ";", F_IMMEDIATE
  mov ax, EXIT
  call COMMA
  jmp short LBRACK
```

The way these words jump to another word at the end is quite convenient. Remember how the NEXT s are written out as part of the defcode of the next word? One of the words needs to be last in memory, and then it won't have any "next word" after it. : and ; are perfect candidates for this, since they don't need a NEXT at all.

# Loading code from disk

Since we don't want to type in disk routines on every boot, we need to include a way to run source code loaded from disk. A filesystem would be its own beast, but Forth tradition has a minimalistic solution: the disk is simply divided into 1 KiB blocks, in which source code is stored, formatted as 16 lines of 64 characters. Then load (blknum -- ) will execute the block with the specified number.

We map block 0 into LBA 0 and 1, block 1 into LBA 2 and 3, and so on. This does mean that block 0 is partially taken by the MBR and LBA 1 is wasted, but I'm not particularily bothered by that.

Since the original BIOS service at int 0x13 / ah = 0x02 requires CHS addressing, I decided to use the EDD extension variant (ah = 0x42). This does mean that floppies are not supported, but I wasn't planning on using any anyway.

To use the EDD interface, we need to build a disk address packet, which looks like this:

```
db 0x10 ; size of packet
db 0  ; reserved
dw sector_count
dw buffer_offset, buffer_segment
dq LBA
```

We use a hybrid strategy to create this packet. The first part is kept as data in the bootsector, but the rest is written at runtime, even if it doesn't change. The "template" needs to be in a place where we can write after it, so the perfect place is just before the compressed data:

```
DiskPacket:
    db 0x10, 0
.count:
    dw 2
.buffer:
    ; rest is filled out at runtime, overwriting the compressed data,
    ; which isn't necessary anymore

CompressedData:
    times COMPRESSED SIZE db 0xcc
```

The first four bytes of the packet are random enough to be hardcoded. However, when it comes to the address of the buffer, we can do better. We will need to write said address to <code>InputPtr</code> anyway. The most direct way to do that takes six bytes:

```
c706020a0006 mov word[InputPtr], BlockBuf
```

However, we can get that value in AX at no extra cost:

```
b80006 mov ax, BlockBuf a3020a mov [InputPtr], ax
```

Thus, we can write these two bytes of the disk packet with only 1 byte of code:

```
defcode LOAD, "load"
   pusha
   mov di, DiskPacket.buffer
   mov ax, BlockBuf
   mov word[InputPtr], ax
   stosw
```

Fitting a Forth in 512 bytes

Next, we need to write the segment ( 0000 ) and the LBA (which ends in six 00 bytes). I like to think of the instructions corresponding to these like so:

```
31c0xor ax, ax; LBA zeroesabstosw; segmentdle3shl bx, 1; LBA data93xchg ax, bx; LBA dataabstosw; LBA data93xchg ax, bx; segmentabstosw; LBA zeroesabstosw; LBA zeroesabstosw; LBA zeroes
```

That is, we write the six LBA zeroes in 5 bytes of code. Writing out the segment only took moving the xor ax, ax earlier, and an additional stosw and xchg ax, bx. Thus, it is neutral at 2 bytes (but we need to write it out in code so that the pointer is right for the rest of the packet). Lastly, of course, we have the actual LBA data, which changes.

While AX is zero, let's take this opportunity to poke in a null terminator after the buffer:

```
mov [BlockBuf.end], al
```

Now we're ready to call the BIOS function. If it errors out, we just loop, as recovering is complicated — the most annoying complication is that the sector count in the packet is overwritten by the number of sectors *successfully read*, which breaks our template.

```
DRIVE_NUMBER equ $+1

mov d1, 0

mov ah, 0x42

mov si, DiskPacket

int 0x13

jc short $

popa

pop bx
```

# Printing numbers

u. prints an unsigned number, followed by a space. Since splitting the number into digits with division yields the least-significant digit first, we push the digits onto the stack, and then pop and print in a separate loop. The space is printed by pushing a fake "digit" that will get converted into a space. This also lets us detect when we popped all the digits — the printing loop stops when it just printed a space.

```
defcode UDOT, "u."
    xchg ax, bx
    push " " - "0"
.split:
   xor dx, dx
   div word[BASE]
   push dx
    or ax, ax
   jnz .split
.print:
   pop ax
   add al, "0"
   cmp al, "9"
   jbe .got digit
   add al, "A" - "0" - 10
.got digit:
    call PutChar
    cmp al, " "
    jne short .print
    pop bx
```

## s: — string poke

s: is a feature which is, I believe, uniquely relevant to bootstrapping. This word takes the address of a buffer, and copies the rest of the current input line there. Without this, a significant amount of code would've had to be typed in twice: first to actually run it and bootstrap a disk block editor, and then again to actually save it on disk.

The implementation is just a simple loop, but the setup around it is noteworthy — we want to load the input pointer into si, but we also need to preserve si so that we can return properly. By using xchg, we can preserve it in [inputPtr] for the duration of the copy, at no extra cost:

```
;; Copies the rest of the line to buf.
defcode LINE, "s:" ; ( buf -- buf+len )
```

```
xchg si, [InputPtr]
.copy:
   lodsb
   mov [bx], al
   inc bx
   or al, al
   jnz short .copy
.done:
   dec bx
   dec si
   xchg si, [InputPtr]
```

The destination pointer is kept in BX. While writing at DI would only take a stosb, getting the pointer in and out of DI outweights this benefit. At the end, we leave a pointer to the null terminator on the stack. That way, you can continue the string by just using S: again on the next line. Since we don't skip any leading whitespace, this is even guaranteed to be properly spaced.

### Other primitives

Choosing the primitives to include in Miniforth is perhaps the biggest tradeoff to be made. I am fully expecting that some more primitive words will need to be defined at runtime by poking in opcodes. After all, there aren't any branching words. However, I'm pretty certain that these opcodes will be able to be generated by a simple Forth assembler, rather than simply hardcoded.

Arithmetic as basic as + is indispensible. I am defining both + and -, though, if I wanted to fit in something more important, I could keep only - and later define : negate 0 swap -; and : + negate -; .

Like any low-level programming language, we need a way to peek and poke values into memory. The implementation of ! is particularly nice, since we can just pop directly into <code>[bx]</code>:

```
defcode PEEK, "@" ; ( addr -- val )
    mov bx, [bx]

defcode POKE, "!" ; ( val addr -- )
    pop word [bx]
    pop bx
```

There also are variants that read and write only a single byte:

```
defcode CPEEK, "c@" ; ( addr -- ch )
    movzx bx, byte[bx]
```

```
defcode CPOKE, "c!" ; ( ch addr -- )
    pop ax
    mov [bx], al
    pop bx
```

We certainly need some stack manipulation words. dup and drop have dead simple implementations, and swap is definitely too useful to skip it.

```
defcode DUP, "dup" ; ( a -- a a )
    push bx

defcode DROP, "drop" ; ( a -- )
    pop bx

defcode SWAP, "swap" ; ( a b -- b a )
    pop ax
    push bx
    xchg ax, bx
```

I chose to also include >r and r>, which allow using the return stack as a second stack for values (but, obviously, only within a single word). This is quite powerful. In fact, combined with dup, drop and swap, they allow you to implement any stack manipulation word you can imagine. 9

```
defcode TO_R, ">r"
    xchg ax, bx
    stosw
    pop bx

defcode FROM_R, "r>"
    dec di
    dec di
    push bx
    mov bx, [di]
```

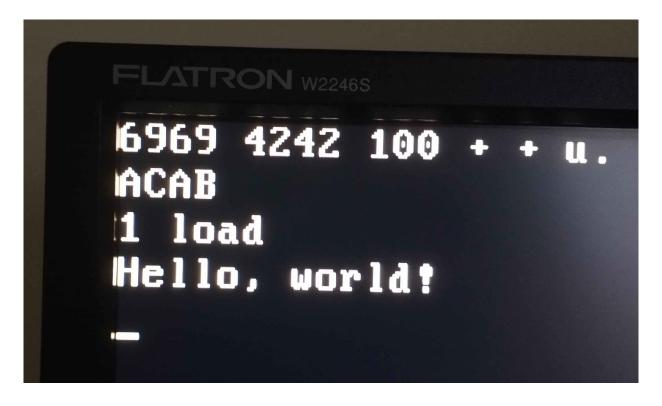
Finally, emit prints a character. This is far enough from the critical path of the bootstrap, that I would be comfortable with removing this one if need be.

```
defcode EMIT, "emit"
    xchg bx, ax
    call PutChar
    pop bx
```

### Conclusion

I am pleased with how this turned out. For a system constrained to the boot sector, I can pretty much call it feature-complete — I can't think of anything that would significantly simplify the bootstrap, while taking few enough bytes that it seems remotely within the reach of code golf. This is largely thanks to Ilya Kurdyukov's help — without it, I wouldn't have been able to fit s: in.

I've found an old PC I can use for my experiments. It boots Miniforth just fine, so I'll be using it from now on. 10



If you'd like to see how to bootstrap on top of Miniforth's minimal set of primitives, see the <u>next post</u> in this series.

### Enjoyed this article?

You might like my <u>other posts</u>, too. If you'd like to be notified of new ones, you can <u>follow me</u> <u>on Twitter</u> or subscribe to the <u>RSS feed</u>.

Thanks to Michał Sidor for supporting my work. Wanna join them?

<sup>1</sup> A graph theorist would have many strong words to describe this, rather than just cycle.  $\stackrel{\frown}{=}$ 

2

And even if that wasn't the case, there are <u>many</u>, <u>many examples</u> of x86 code written with the printable subset of ASCII. I've even <u>done it myself</u> once a few years ago. <u>←</u>

- <sup>3</sup> If you, dear reader, find this unsatisfactory, I would like to invite you on your own journey of bootstrapping. It's really quite fun! <u>←</u>
- <sup>4</sup> The exact savings at the moment are somewhat hard to calculate, because some words don't make use of the NEXT being appended. <u>••</u>
- <sup>5</sup> While this passage of the reference only talks about pop ss, an analogous statement is made in the documentation for mov . =
- 6 This seems to be one of the many mistakes in the SDM using a pop esp for this wouldn't work. Section 6.8.3 ("Masking Exceptions and Interrupts When Switching Stacks") in Volume 3A clarifies that all single-instruction ways to load sp work for this. I would've quoted that section instead, if not for the fact that, while it lists many more types of events that are suppressed, it actually fails to mention actual interrupts as one of them. That section does mention some interesting edge-cases, though. For example, if you're like me, you might be wondering what happens if many instructions in a row write to ss. The answer is that only the first one is guaranteed to suppress interrupts.  $\underline{\leftarrow}$
- <sup>7</sup> I suppose it makes sense that this option is only available for <code>imul</code>, since, by modular arithmetic, the only difference between a signed and unsigned multiply is in the upper half of the product, which we are discarding here. Immediates could be useful with <code>mul</code> too, though...  $\underline{\leftarrow}$
- 8 You might ask, why not just declare that BASE is a byte-sized variable? The answer is that u., which is the word that prints a number, uses div word[BASE], so that the result is still 16-bit.
- 9 This does not include words like PICK you would need loops for that. Anything definable as ( of names> -- of names> ) is fair game, though. Proving this fact is left as an exercise to the reader.  $^{11} \leftarrow$
- 10 A note from the future: I've since switched to dual-booting it on my laptop.
- 11 Seriously though, try to devise an infallible strategy for turning a stack effect into a sequence of words that implement it. It's a nice problem. —

#### Next post:

No branches? No problem — a Forth assembler

**10 Comments** - powered by utteranc.es

NieDzejkob commented on Jun 12, 2021

Owner

Lust making sure that the commenting works :D