

B.Y.O.ASSEMBLER

-or-

Build Your Own (Cross-) Assembler....in Forth

by Brad Rodriguez

A. INTRODUCTION

In a previous issue of this journal I described how to "bootstrap" yourself into a new processor, with a simple debug monitor. But how do you write code for this new CPU, when you can't find or can't afford an assembler? Build your own!

Forth is an ideal language for this. I've written cross-assemblers in as little as two hours (for the TMS320, over a long lunch break). Two days is perhaps more common; and one processor (the Zilog Super8) took me five days. But when you have more time than money, this is a bargain.

In part 1 of this article I will describe the basic principles of Forth-style assemblers -- structured, single-pass, postfix. Much of this will apply to any processor, and these concepts are in almost every Forth assembler.

In part 2 I will examine an assembler for a specific CPU: the Motorola 6809. This assembler is simple but not trivial, occupying 15 screens of source code. Among other things, it shows how to handle instructions with multiple modes (in this case, addressing modes). By studying this example, you can figure out how to handle the peculiarities of your own CPU.

B. WHY USE FORTH?

I believe that Forth is the easiest language in which to write assemblers.

First and foremost, Forth has a "text interpreter" designed to look up text strings and perform some related action. Turning text strings into bytes is exactly what is needed to compile assembler mnemonics! Operands and addressing modes can also be handled as Forth "words."

Forth also includes "defining words," which create large sets of words with a common action. This feature is very useful when defining assembler mnemonics.

Since every Forth word is always available, Forth's arithmetic and logical functions can be used within the assembler environment to perform address and operand arithmetic.

Finally, since the assembler is entirely implemented in Forth words, Forth's "colon definitions" provide a rudimentary macro facility, with no extra effort.

C. THE SIMPLEST CASE: ASSEMBLING A NOP

To understand how Forth translates mnemonics to machine code, consider the simplest case: the NOP instruction (12 hex on the 6809).

A conventional assembler, on encountering a NOP in the opcode field, must append a 12H byte to the output file and advance the location counter by 1. Operands and comments are ignored. (I will ignore labels for the time being.)

In Forth, the memory-resident dictionary is usually the output "file." So, make NOP a Forth word, and give it an action, namely, "append 12H to the dictionary and advance

the dictionary pointer."

```

    HEX
    : NOP,    12 C, ;

```

Assembler opcodes are often given Forth names which include a trailing comma, as shown above. This is because many Forth words -- such as AND XOR and OR -- conflict with assembler mnemonics. The simplest solution is to change the assembler mnemonics slightly, usually with a trailing comma. (This comma is a Forth convention, indicating that something is appended to the dictionary.)

D. THE CLASS OF "INHERENT" OPCODES

Most processors have many instructions, like NOP, which require no operands. All of these could be defined as Forth colon definitions, but this duplicates code, and wastes a lot of space. It's much more efficient to use Forth's "defining word" mechanism to give all of these words a common action. In object-oriented parlance, this builds "instances" of a single "class."

This is done with Forth's CREATE and DOES>. (In fig-Forth, as used in the 6809 assembler, the words are <BUILDS and DOES>.)

```

: INHERENT      ( Defines the name of the class)
  CREATE       ( this will create an instance)
    C,         ( store the parameter for each
                instance)
  DOES>        ( this is the class' common action)
    C@         ( get each instance's parameter)
    C,         ( the assembly action, as above)
  ;           ( End of definition)

HEX
12 INHERENT NOP, ( Defines an instance NOP, of class
                  INHERENT, with parameter 12H.)
3A INHERENT ABX, ( Another instance - the ABX instr)
3D INHERENT MUL, ( Another instance - the MUL instr)

```

In this case, the parameter (which is specific to each instance) is simply the opcode to be assembled for each instruction.

This technique provides a substantial memory savings, with almost no speed penalty. But the real advantage becomes evident when complex instruction actions -- such as required for parameters, or addressing modes -- are involved.

E. HANDLING OPERANDS

Most assembler opcodes, it is true, require one or more operands. As part of the action for these instructions, Forth routines could be written to parse text from the input stream, and interpret this text as operand fields. But why? The Forth environment already provides a parse-and-interpret mechanism!

So, Forth will be used to parse operands. Numbers are parsed normally (in any base!), and equates can be Forth CONSTANTS. But, since the operands determine how the opcode is handled, they will be processed first. The results of operand parsing will be left on Forth's stack, to be picked up by the opcode word. This leads to Forth's unique postfix format for assemblers: operands, followed by opcode.

Take, for example, the 6809's ORCC instruction, which takes a single numeric parameter:

```

    HEX
    : ORCC,   1A C, C, ;

```

The exact sequence of actions for ORCC, is: 1) put 1A hex on the parameter stack; 2) append the top stack item (the 1A) to the dictionary, and drop it from the stack; 3) append the new top stack item (the operand) to the dictionary, and drop it from the stack. It is assumed that a numeric value was already on the stack, for the second C, to use. This numeric value is the result of the operand parsing, which, in this case, is simply the parsing of a single integer value:

```

    HEX
    0F ORCC,

```

The advantage here is that all of Forth's power to operate on stack values, via both built-in operators and newly-defined functions, can be employed to create and modify operands. For example:

```

    HEX
    01 CONSTANT CY-FLAG      ( a "named" numeric value)
    02 CONSTANT OV-FLAG
    04 CONSTANT Z-FLAG
    ...
    CY-FLAG Z-FLAG + ORCC,   ( add 1 and 4 to get operand)

```

The extension of operand-passing to the defining words technique is straightforward.

F. HANDLING ADDRESSING MODES

Rarely can an operand, or an opcode, be used unmodified. Most of the instructions in a modern processor can take multiple forms, depending on the programmer's choice of addressing mode.

Forth assemblers have attacked this problem in a number of ways, depending on the requirements of the specific processor. All of these techniques remain true to the Forth methodology: the addressing mode operators are implemented as Forth words. When these words are executed, they alter the assembly of the current instruction.

1. Leaving additional parameters on the stack.

This is most useful when an addressing mode must always be specified. The addressing-mode word leaves some constant value on the stack, to be picked up by the opcode word. Sometimes this value can be a "magic number" which can be added to the opcode to modify it for the different mode. When this is not feasible, the addressing-mode value can activate a CASE statement within the opcode, to select one of several actions. In this latter case, instructions of different lengths, possibly with different operands, can be assembled depending on the addressing mode.

2. Setting flags or values in fixed variables.

This is most useful when the addressing mode is optional. Without knowing whether an addressing mode was specified, you don't know if the value on the stack is a "magic number" or just an operand value. The solution: have the addressing mode put its magic number in a predefined variable (often called MODE). This variable is initialized to a default value, and reset to this default value after each instruction is assembled. Thus, this variable can be tested to see if an addressing mode was specified (overriding the default).

3. Modifying parameter values already on the stack.

It is occasionally possible to implement addressing mode words that work by modifying an operand value. This is rarely seen.

All three of these techniques are used, to some extent, within the 6809 assembler.

For most processors, register names can simply be Forth CONSTANTS, which leave a value on the stack. For some processors it is useful to have register names specify "register addressing mode" as well. This is easily done by defining register names with a new defining word, whose run-time action sets the addressing mode (either on the stack or in a MODE variable).

Some processors allow multiple addressing modes in a single instruction. If the number of addressing modes is fixed by the instruction, they can be left on the stack. If the number of addressing modes is variable, and it is desired to know how many have been specified, multiple MODE variables can be used for the first, second, etc. (In one case -- the Super8 -- I had to keep track of not only how many addressing modes were specified, but also where among the operands they were specified. I did this by saving the stack position along with each addressing mode.)

Consider the 6809 ADD instruction. To simplify things, ignore the Indexed addressing modes for now, and just consider the remaining three addressing modes: Immediate, Direct, and Extended. These will be specified as follows:

	source code	assembles as
Immediate:	number # ADD,	8B nn
Direct:	address <> ADD,	9B aa
Extended:	address ADD,	BB aa aa

Since Extended has no addressing mode operator, the mode-variable approach seems to be indicated. The Forth words # and <> will set MODE.

Observe the regularity in the 6809 opcodes. If the Immediate opcode is the "base" value, then the Direct opcode is this value plus 10 hex, and the Extended opcode is this value plus 30 hex. (And the Indexed opcode, incidentally, is this value plus 20 hex.) This applies uniformly across almost all 6809 instructions which use these addressing modes. (The exceptions are those opcodes whose Direct opcodes are of the form 0x hex.)

Regularities like this are made to be exploited! This is a general rule for writing assemblers: find or make an opcode chart, and look for regularities -- especially those applying to addressing modes or other instruction modifiers (like condition codes).

In this case, appropriate MODE values are suggested:

```
VARIABLE MODE HEX
: #      0 MODE ! ;
: <>    10 MODE ! ;
: RESET  30 MODE ! ;
```

The default MODE value is 30 hex (for Extended mode), so a Forth word RESET is added to restore this value. RESET will be used after every instruction is assembled.

The ADD, routine can now be written. Let's go ahead and write it using a defining word:

```
HEX
: GENERAL-OP \ base-opcode --
  CREATE C,
  DOES> \ operand --
  C@ \ get the base opcode
  MODE @ + \ add the "magic number"
  C, \ assemble the opcode
  MODE @ CASE
    0 OF C, ENDOF \ byte operand
```

```

10 OF C, ENDOF    \ byte operand
30 OF , ENDOF     \ word operand
ENDCASE
RESET ;

```

8B GENERAL-OP ADD,

Each "instance" of GENERAL-OP will have a different base opcode. When ADD, executes, it will fetch this base opcode, add the MODE value to it, and assemble that byte. Then it will take the operand which was passed on the stack, and assemble it either as a byte or word operand, depending on the selected mode. Finally, it will reset MODE.

Note that all of the code is now defined to create instructions in the same family as ADD:

```

HEX 89 GENERAL-OP ADC,
    84 GENERAL-OP AND,
    85 GENERAL-OP BIT,
    etc.

```

The memory savings from defining words really become evident now. Each new opcode word executes the lengthy bit of DOES> code given above; but each word is only a one-byte Forth definition (plus header and code field, of course).

This is not the actual code from the 6809 assembler -- there are additional special cases which need to be handled. But it demonstrates that, by storing enough mode information, and by making liberal use of CASE statements, the most ludicrous instruction sets can be assembled.

G. HANDLING CONTROL STRUCTURES

The virtues of structured programming, have long been sung -- and there are countless "structured assembly" macro packages for conventional assemblers. But Forth assemblers favor label-free, structured assembly code for a pragmatic reason: in Forth, it's simpler to create assembler structures than labels!

The structures commonly included in Forth assemblers are intended to resemble the programming structures of high-level Forth. (Again, the assembler structures are usually distinguished by a trailing comma.)

1. BEGIN, ... UNTIL,

The BEGIN, ... UNTIL, construct is the simplest assembler structure to understand. The assembler code is to loop back to the BEGIN point, until some condition is satisfied. The Forth assembler syntax is

```
BEGIN,    more code    cc UNTIL,
```

where 'cc' is a condition code, which has presumably been defined -- either as an operand or an addressing mode -- for the jump instructions.

Obviously, the UNTIL, will assemble a conditional jump. The sense of the jump must be "inverted" so that if 'cc' is satisfied, the jump does NOT take place, but instead the code "falls through" the jump. The conventional assembler equivalent would be:

```

xxx:  ...
      ...
      ...
      JR  ~cc,xxx

```

(where ~cc is the logical inverse of cc.)

Forth offers two aids to implementing BEGIN, and UNTIL,. The word HERE will return the current location counter value. And values may be kept deep in the stack, with no effect on Forth processing, then "elevated" when required.

So: BEGIN, will "remember" a location counter, by placing its value on the stack. UNTIL, will assemble a conditional jump to the "remembered" location.

```
: BEGIN, ( - a)      HERE ;
: UNTIL, ( a cc - ) NOTCC JR, ;
```

This introduces the common Forth stack notation, to indicate that BEGIN, leaves one value (an address) on the stack. UNTIL, consumes two values (an address and a condition code) from the stack, with the condition code on top. It is presumed that a word NOTCC has been defined, which will convert a condition code to its logical inverse. It is also presumed that the opcode word JR, has been defined, which will expect an address and a condition code as operands. (JR, is a more general example than the branch instructions used in the 6809 assembler.)

The use of the stack for storage of the loop address allows BEGIN, ... UNTIL, constructs to be nested, as:

```
BEGIN, ... BEGIN, ... cc UNTIL, ... cc UNTIL,
```

The "inner" UNTIL, resolves the "inner" BEGIN, forming a loop wholly contained within the outer BEGIN, ... UNTIL, loop.

2. BEGIN, ... AGAIN,

Forth commonly provides an "infinite loop" construct, BEGIN ... AGAIN, which never exits. For the sake of completeness, this is usually implemented in the assembler as well.

Obviously, this is implemented in the same manner as BEGIN, ... UNTIL, except that the jump which is assembled by AGAIN, is an unconditional jump.

3. DO, ... LOOP,

Many processors offer some kind of looping instruction. Since the 6809 does not, let's consider the Zilog Super8; its Decrement-and-Jump-Non-Zero (DJNZ) instruction can use any of 16 registers as the loop counter. This can be written in structured assembler:

```
DO,      more code    r LOOP,
```

where r is the register used as the loop counter. Once again, the intent is to make the assembler construct resemble the high-level Forth construct.

```
: DO,    ( - a)      HERE ;
: LOOP,  ( a r - )    DJNZ, ;
```

Some Forth assemblers go so far as to make DO, assemble a load-immediate instruction for the loop counter -- but this loses flexibility. Sometimes the loop count isn't a constant. So I prefer the above definition of DO, .

4. IF, ... THEN,

The IF, ... THEN, construct is the simplest forward-referencing construct. If a condition is satisfied, the code within the IF,...THEN, is to be executed; otherwise, control is transferred to the first instruction after THEN,.

(Note that Forth normally employs THEN, where other languages use "endif." You can have both in your

assembler.)

The Forth syntax is

```
cc IF, ... .. THEN,
```

for which the "conventional" equivalent is

```
JP ~cc,xxx
...
...
...
xxx:
```

Note that, once again, the condition code must be inverted to produce the expected logical sense for IF, .

In a single pass assembler, the requisite forward jump cannot be directly assembled, since the destination address of the jump is not known when IF, is encountered. This problem is solved by causing IF, to assemble a "dummy" jump, and stack the address of the jump's operand field. Later, the word THEN, (which will provide the destination address) can remove this stacked address and "patch" the jump instruction accordingly.

```
: IF, ( cc - a) NOT 0 SWAP JP, ( conditional jump
      HERE 2 - ; with 2-byte operand)
: THEN, ( a) HERE SWAP ! ; ( store HERE at the
                          stacked address)
```

IF, inverts the condition code, assembles a conditional jump to address zero, and then puts on the stack the address of the jump address field. (After JP, is assembled, the location counter HERE points past the jump instruction, so we need to subtract two to get the location of the address field.) THEN, will patch the current location into the operand field of that jump.

If relative jumps are used, additional code must be added to THEN, to calculate the relative offset.

5. IF, ... ELSE, ... THEN,

A refinement of the IF,...THEN, construct allows code to be executed if the condition is NOT satisfied. The Forth syntax is

```
cc IF, ... .. ELSE, ... .. THEN,
```

ELSE, has the expected meaning: if the first part of this statement is not executed, then the second part is.

The assembler code necessary to create this construct is:

```
JP ~cc,xxx
... ( the "if" code)
...
JP yyy
xxx: ... ( the "else" code)
...
yyy:
```

ELSE, must modify the actions of IF, and THEN, as follows: a) the forward jump from IF, must be patched to the start of the "else" code ("xxx"); and b) the address supplied by THEN, must be patched into the unconditional jump instruction at the end of the "if" code ("JP yyy"). ELSE, must also assemble the unconditional jump. This is done thus:

```
: ELSE ( a - a) 0 T JP, ( unconditional jump)
      HERE 2 - ( stack its address
               for THEN, to patch)
```

```

      SWAP      ( get the patch address
                  of the IF, jump)
      HERE SWAP ! ( patch it to the current
                  location, i.e., the
                  next instruction)
      ;

```

Note that the jump condition 'T' assembles a "jump always" instruction. The code from IF, and THEN, can be "re-used" if the condition 'F' is defined as the condition-code inverse of 'T':

```

      : ELSE ( a - a)  F IF,  SWAP  THEN, ;

```

The SWAP of the stacked addresses reverses the patch order, so that the THEN, inside ELSE, patches the original IF; and the final THEN, patches the IF, inside ELSE,. Graphically, this becomes:

```

      IF,(1)  ...  IF,(2) THEN,(1)  ...  THEN,(2)
                  \_____/
                  inside ELSE,

```

IF,...THEN, and IF,...ELSE,...THEN, structures can be nested. This freedom of nesting also extends to mixtures of these and BEGIN,...UNTIL, structures.

6. BEGIN, ... WHILE, ... REPEAT,

The final, and most complex, assembler control structure is the "while" loop in which the condition is tested at the beginning of the loop, rather than at the end.

In Forth the accepted syntax for this structure is

```

      BEGIN, evaluate cc WHILE, loop code REPEAT,

```

In practice, any code -- not just condition evaluations -- may be inserted between BEGIN, and WHILE,.

What needs to be assembled is this: WHILE, will assemble a conditional jump, on the inverse of cc, to the code following the REPEAT,. (If the condition code cc is satisfied, we should "fall through" WHILE, to execute the loop code.) REPEAT, will assemble an unconditional jump back to BEGIN. Or, in terms of existing constructs:

```

      BEGIN,(1) ... cc IF,(2) ... AGAIN,(1) THEN,(2)

```

Once again, this can be implemented with existing words, by means of a stack manipulation inside WHILE, to re-arrange what jumps are patched by whom:

```

      : WHILE, ( a cc - a a)  IF, SWAP ;
      : REPEAT, ( a a - )  AGAIN, THEN, ;

```

Again, nesting is freely permitted.

H. THE FORTH DEFINITION HEADER

In most applications, machine code created by a Forth assembler will be put in a CODE word in the Forth dictionary. This requires giving it an identifying text "name," and linking it into the dictionary list.

The Forth word CREATE performs these functions for the programmer. CREATE will parse a word from the input stream, build a new entry in the dictionary with that name, and adjust the dictionary pointer to the start of the "definition field" for this word.

Standard Forth uses the word CODE to distinguish the start of an assembler definition in the Forth dictionary. In addition to performing CREATE, the word CODE may set the

assembler environment (vocabulary), and may reset variables (such as MODE) in the assembler. Some Forths may also require a "code address" field; this is set by CREATE in some systems, while others expect CODE to do this.

I. SPECIAL CASES

1. Resident vs. cross-compilation

Up to now, it has been assumed that the machine code is to be assembled into the dictionary of the machine running the assembler.

For cross-assembly and cross-compilation, code is usually assembled for the "target" machine into a different area of memory. This area may or may not have its own dictionary structure, but it is separate from the "host" machine's dictionary.

The most common and straightforward solution is to provide the host machine with a set of Forth operators to access the "target" memory space. These are made deliberately analogous to the normal Forth memory and dictionary operators, and are usually distinguished by the prefix "T". The basic set of operators required is:

TDP	target dictionary pointer DP
THERE	analogous to HERE, returns TDP
TC,	target byte append C,
TC@	target byte fetch C@
TC!	target byte store C!
T@	target word fetch @
T!	target word store !

Sometimes, instead of using the "T" prefix, these words will be given identical names but in a different Forth vocabulary. (The vocabulary structure in Forth allows unambiguous use of the same word name in multiple contexts.) The 6809 assembler in Part 2 assumes this.

2. Compiling to disk

Assembler output can be directed to disk, rather than to memory. This, too, can be handled by defining a new set of dictionary, fetch, and store operators. They can be distinguished with a different prefix (such as "T" again), or put in a distinct vocabulary.

Note that the "patching" manipulations used in the single-pass control structures require a randomly-accessible output medium. This is not a problem with disk, although heavy use of control structures may result in some inefficient disk access.

3. Compiler Security

Some Forth implementations include a feature known as "compiler security," which attempts to catch mismatches of control structures. For example, the structure

```
IF, ... cc UNTIL,
```

would leave the stack balanced (UNTIL, consumes the address left by IF,), but would result in nonsense code.

The usual method for checking the match of control structures is to require the "leading" control word to leave a code value on the stack, and the "trailing" word to check the stack for the correct value. For example:

```
IF, leaves a 1;
THEN, checks for a 1;
ELSE, checks for a 1 and leaves a 1;
```

```

BEGIN,  leaves a 2;
UNTIL,  checks for a 2;
AGAIN,  checks for a 2;
WHILE,  checks for a 2 and leaves a 3;
REPEAT,  checks for a 3.

```

This will detect most mismatches. Additional checks may be included for the stack imbalance caused by "unmatched" control words. (The 6809 assembler uses both of these error checks.)

The cost of compiler security is the increased complexity of the stack manipulations in such words as ELSE, and WHILE,. Also, the programmer may wish to alter the order in which control structures are resolved, by manually re-arranging the stack; compiler security makes this more difficult.

4. Labels

Even in the era of structured programming, some programmers will insist on labels in their assembler code.

The principal problem with named labels in a Forth assembler definition is that the labels themselves are Forth words. They are compiled into the dictionary -- usually at an inconvenient point, such as inside the machine code. For example:

```

CODE TEST ... machine code ...
  HERE CONSTANT LABEL1
  ... machine code ...
  LABEL1 NZ JP,

```

will cause the dictionary header for LABEL1 -- text, links, and all -- to be inserted in the middle of CODE. Several solutions have been proposed:

- a) define labels only "outside" machine code.
Occasionally useful, but very restricted.
- b) use some predefined storage locations (variables) to provide "temporary," or local, labels.
- c) use a separate dictionary space for the labels, e.g., as provided by the TRANSIENT scheme [3].
- d) use a separate dictionary space for the machine code.
This is common practice for meta-compilation; most Forth meta- compilers support labels with little difficulty.

5. Table Driven Assemblers

Most Forth assemblers can handle the profusion of addressing modes and instruction opcodes by CASE statements and other flow-of-control constructs. These may be referred to as "procedural" assemblers.

Some processors, notably the Motorola 68000, have instruction and addressing sets so complex as to render the decision trees immense. In such cases, a more "table-driven" approach may save substantial memory and processor time.

(I avoid such processors. Table driven assemblers are much more complex to write.)

6. Prefix Assemblers

Sometimes a prefix assembler is unavoidable. (One example: I recently translated many K of Super8 assembler code from the Zilog assembler to a Forth assembler.) There is a programming "trick" which simulates a prefix assembler, while using the assembler techniques described in this

article.

Basically, this trick is to "postpone" execution of the opcode word, until after the operands have been evaluated. How can the assembler determine when the operands are finished? Easy: when the next opcode word is encountered.

So, every opcode word is modified to a) save its own execution address somewhere, and b) execute the "saved" action of the previous opcode word. For example:

```
... JP operand    ADD operands ...
```

JP stores its execution address (and the address of its "instance" parameters) in a variable somewhere. Then, the operands are evaluated. ADD will fetch the information saved by JP, and execute the run-time action of JP. The JP action will pick up whatever the operands left on the stack. When the JP action returns, ADD will save its own execution address and instance parameters, and the process continues. (Of course, JP would have executed its previous opcode.)

This is confusing. Special care must be taken for the first and last opcodes in the assembler code. If mode variables are used, the problem of properly saving and restoring them becomes nightmarish. I leave this subject as an exercise for the advanced student...or for an article of its own.

J. CONCLUSION

I've touched upon the common techniques used in Forth assemblers. Since I believe the second-best way to learn is by example, in part 2 I will present the full code for the 6809 assembler. Studying a working assembler may give you hints on writing an assembler of your own.

The BEST way to learn is by doing!

K. REFERENCES

1. Curley, Charles, Advancing Forth. Unpublished manuscript (1985).
2. Wasson, Philip, "Transient Definitions," Forth Dimensions III/6 (Mar-Apr 1982), p.171.

L. ADDITIONAL SOURCES

1. Cassady, John J., "8080 Assembler," Forth Dimensions III/6 (Mar-Apr 1982), pp. 180-181. Noteworthy in that the entire assembler fits in less than 48 lines of code.
2. Ragsdale, William F., "A FORTH Assembler for the 6502," Dr. Dobb's Journal #59 (September 1981), pp. 12-24. A simple illustration of addressing modes.
3. Duncan, Ray, "FORTH 8086 Assembler," Dr. Dobb's Journal #64 (February 1982), pp. 14-18 and 33-46.
4. Perry, Michael A., "A 68000 Forth Assembler," Dr. Dobb's Journal #83 (September 1983), pp. 28-42.
5. Assemblers for the 8080, 8051, 6502, 68HC11, 8086, 80386, 68000, SC32, and Transputer can be downloaded from the Forth Interest Group (FORTH) conference on GENie.

□